

DQML Final Project
Video_Game_Sales_and_Analytics
Database

**Performance Tuning Report On Complex
Queries**

1. Introduction

This report documents the performance tuning process for the analytical queries created in Phase 2 (OLAP Layer). Based on EXPLAIN ANALYZE results, the most expensive query was selected and optimized using a targeted indexing strategy.

The goal of this step was to:

- Identify the slowest analytical query
- Analyze its execution plan
- Apply appropriate indexes
- Measure performance before and after optimization
- Explain why the indexing strategy improved performance

2. Identifying the Most Complex Query

All five analytical queries from Phase 2 were tested using:

EXPLAIN ANALYZE <query>

QUERY 1: BEST PERFORMING GENRE FOR EACH PLATFORM

```
1  /* QUERY 1: BEST PERFORMING GENRE FOR EACH PLATFORM */
2
3  EXPLAIN ANALYZE
4  WITH genre_sales AS (
5      SELECT
6          p.platform_name,
7          g.genre_id,
8          ge.genre_name,
9          SUM(s.global_sales) AS total_sales
10     FROM game g
11     JOIN platform p ON g.platform_id = p.platform_id
12     JOIN genre ge ON g.genre_id = ge.genre_id
13     JOIN sales s ON g.game_id = s.game_id
14     GROUP BY p.platform_name, g.genre_id, ge.genre_name
15 ), ranked AS (
16     SELECT *,
17         RANK() OVER (PARTITION BY platform_name ORDER BY total_sales DESC) AS rnk
18     FROM genre_sales
19 )
20
21     SELECT platform_name, genre_name, total_sales
22     FROM ranked
23     WHERE rnk = 1;
```

Output:

The screenshot shows a database interface with a "QUERY PLAN" tab selected. The plan details the execution steps for a query involving multiple tables (game, publisher, reviews) and various operations like Subquery Scan, WindowAgg, Sort, HashAggregate, Hash Join, Seq Scan, and Hash. The plan includes cost estimates, actual times, and memory usage per step. The total execution time is 15.867 ms.

```
1 Subquery Scan on ranked (cost=2791.06..3332.34 rows=83 width=43) (actual time=15.400..15.445 rows=31 loops=1)
2   Filter: (ranked.rnk = 1)
3     -> WindowAgg (cost=2791.06..3124.16 rows=16655 width=55) (actual time=15.399..15.441 rows=31 loops=1)
4       Run Condition: (rank() OVER () <= 1)
5       -> Sort (cost=2791.06..2832.69 rows=16655 width=43) (actual time=15.377..15.389 rows=294 loops=1)
6         Sort Key: genre_sales.platform_name, genre_sales.total_sales DESC
7         Sort Method: quicksort Memory: 38kB
8         -> Subquery Scan on genre_sales (cost=1248.50..1623.24 rows=16655 width=43) (actual time=14.992..15.177 rows=294 loops=1)
9           -> HashAggregate (cost=1248.50..1456.69 rows=16655 width=47) (actual time=14.991..15.157 rows=294 loops=1)
10          Group Key: p.platform_name, g.genre_id, ge.genre_name
11          Batches: 1 Memory Usage: 913kB
12          -> Hash Join (cost=619.06..1081.95 rows=16655 width=20) (actual time=3.918..11.084 rows=16655 loops=1)
13            Hash Cond: (g.game_id = s.game_id)
14            -> Hash Join (cost=118.33..537.48 rows=16655 width=19) (actual time=0.126..5.148 rows=16655 loops=1)
15              Hash Cond: (g.genre_id = ge.genre_id)
16              -> Hash Join (cost=62.42..437.77 rows=16655 width=11) (actual time=0.068..3.359 rows=16655 loops=1)
17                Hash Cond: (g.platform_id = p.platform_id)
18                -> Seq Scan on game g (cost=0.00..331.55 rows=16655 width=12) (actual time=0.006..0.961 rows=16655 loops=1)
19                -> Hash (cost=33.30..33.30 rows=2330 width=7) (actual time=0.020..0.020 rows=31 loops=1)
20                  Buckets: 4096 Batches: 1 Memory Usage: 34kB
21                  -> Seq Scan on platform p (cost=0.00..33.30 rows=2330 width=7) (actual time=0.011..0.014 rows=31 loops=1)
22                  -> Hash (cost=30.40..30.40 rows=2040 width=12) (actual time=0.048..0.048 rows=13 loops=1)
23                    Buckets: 2048 Batches: 1 Memory Usage: 17kB
24                    -> Seq Scan on genre ge (cost=0.00..30.40 rows=2040 width=12) (actual time=0.032..0.034 rows=13 loops=1)
25                    -> Hash (cost=292.55..292.55 rows=16655 width=9) (actual time=3.704..3.704 rows=16655 loops=1)
26                      Buckets: 32768 Batches: 1 Memory Usage: 972kB
27                      -> Seq Scan on sales s (cost=0.00..292.55 rows=16655 width=9) (actual time=0.005..1.563 rows=16655 loops=1)
28 Planning Time: 1.873 ms
29 Execution Time: 15.867 ms
```

Total rows: 29 Query complete 00:00:00.045

QUERY 2: PUBLISHERS WITH THE BIGGEST CRITIC-USER SCORE GAP

```
26  /* QUERY 2: PUBLISHERS WITH THE BIGGEST CRITIC-USER SCORE GAP */
27
28  EXPLAIN ANALYZE
29  WITH score_gap AS (
30    SELECT |
31      pu.publisher_name,
32      AVG(r.critic_score) AS avg_critic,
33      AVG(r.user_score) AS avg_user,
34      ABS(AVG(r.critic_score) - AVG(r.user_score)) AS gap
35    FROM game g
36    JOIN publisher pu ON g.publisher_id = pu.publisher_id
37    JOIN reviews r ON g.game_id = r.game_id
38    GROUP BY pu.publisher_name
39  )
40  SELECT *,
41        RANK() OVER (ORDER BY gap DESC) AS rank_by_gap
42  FROM score_gap;
```

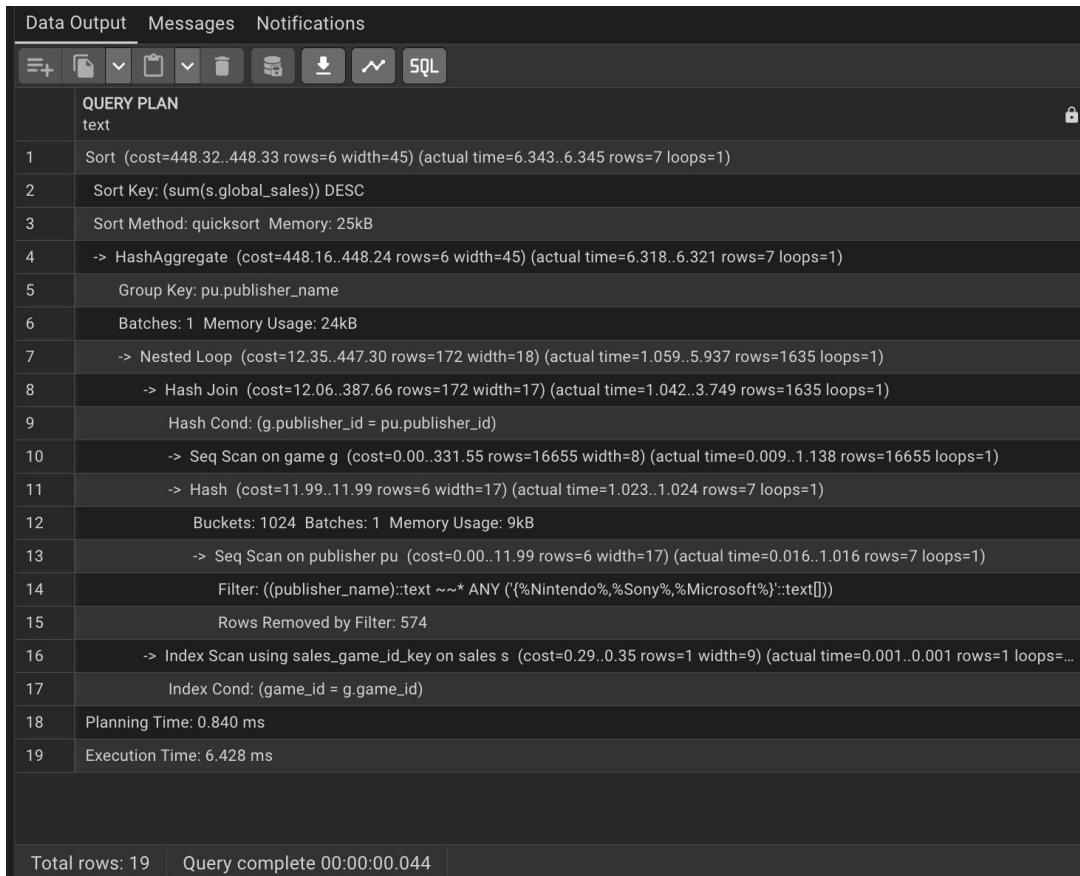
Output:

Data Output		Messages	Notifications
		SQL	
QUERY PLAN			
	text		
1	WindowAgg (cost=1089.34..1099.51 rows=581 width=117) (actual time=15.139..15.316 rows=579 loops=1)		
2	-> Sort (cost=1089.34..1090.79 rows=581 width=109) (actual time=15.133..15.154 rows=579 loops=1)		
3	Sort Key: score_gap.gap DESC		
4	Sort Method: quicksort Memory: 61kB		
5	-> Subquery Scan on score_gap (cost=1051.05..1062.66 rows=581 width=109) (actual time=14.699..15.004 rows=579 loops=1)		
6	-> HashAggregate (cost=1051.05..1062.66 rows=581 width=109) (actual time=14.698..14.962 rows=579 loops=1)		
7	Group Key: pu.publisher_name		
8	Batches: 1 Memory Usage: 553kB		
9	-> Hash Join (cost=506.81..926.13 rows=16655 width=20) (actual time=4.829..10.902 rows=16655 loops=1)		
10	Hash Cond: (g.game_id = r.game_id)		
11	-> Hash Join (cost=17.07..392.66 rows=16655 width=17) (actual time=0.273..3.897 rows=16655 loops=1)		
12	Hash Cond: (g.publisher_id = pu.publisher_id)		
13	-> Seq Scan on game g (cost=0.00..331.55 rows=16655 width=8) (actual time=0.009..1.013 rows=16655 loops=1)		
14	-> Hash (cost=9.81..9.81 rows=581 width=17) (actual time=0.247..0.248 rows=581 loops=1)		
15	Buckets: 1024 Batches: 1 Memory Usage: 37kB		
16	-> Seq Scan on publisher pu (cost=0.00..9.81 rows=581 width=17) (actual time=0.056..0.127 rows=581 loops=1)		
17	-> Hash (cost=281.55..281.55 rows=16655 width=11) (actual time=4.540..4.540 rows=16655 loops=1)		
18	Buckets: 32768 Batches: 1 Memory Usage: 1001kB		
19	-> Seq Scan on reviews r (cost=0.00..281.55 rows=16655 width=11) (actual time=0.006..1.787 rows=16655 loops=1)		
20	Planning Time: 0.629 ms		
21	Execution Time: 15.427 ms		
Total rows: 21		Query complete 00:00:00.041	

QUERY 3: SALES DOMINANCE OF NINTENDO, SONY, MICROSOFT

```
45  /* QUERY 3: SALES DOMINANCE OF NINTENDO, SONY, MICROSOFT */
46
47  EXPLAIN ANALYZE
48  SELECT
49      pu.publisher_name,
50      SUM(s.global_sales) AS total_global_sales
51  FROM game g
52  JOIN publisher pu ON g.publisher_id = pu.publisher_id
53  JOIN sales s ON g.game_id = s.game_id
54  WHERE pu.publisher_name ILIKE ANY (ARRAY['%Nintendo%', '%Sony%', '%Microsoft%'])
55  GROUP BY pu.publisher_name
56  ORDER BY total_global_sales DESC;
```

Output:



The screenshot shows the PostgreSQL Explain Plan interface. The top navigation bar includes tabs for Data Output, Messages, Notifications, and SQL. Below the navigation bar is a toolbar with various icons. The main area displays the query plan text, which details the execution steps, including sorting, aggregation, joins, and scans, along with their respective costs and actual execution times. The plan ends with planning and execution times.

```
1 Sort (cost=448.32..448.33 rows=6 width=45) (actual time=6.343..6.345 rows=7 loops=1)
2   Sort Key: (sum(s.global_sales)) DESC
3   Sort Method: quicksort Memory: 25kB
4     -> HashAggregate (cost=448.16..448.24 rows=6 width=45) (actual time=6.318..6.321 rows=7 loops=1)
5       Group Key: pu.publisher_name
6       Batches: 1 Memory Usage: 24kB
7     -> Nested Loop (cost=12.35..447.30 rows=172 width=18) (actual time=1.059..5.937 rows=1635 loops=1)
8       -> Hash Join (cost=12.06..387.66 rows=172 width=17) (actual time=1.042..3.749 rows=1635 loops=1)
9         Hash Cond: (g.publisher_id = pu.publisher_id)
10        -> Seq Scan on game g (cost=0.00..331.55 rows=16655 width=8) (actual time=0.009..1.138 rows=16655 loops=1)
11        -> Hash (cost=11.99..11.99 rows=6 width=17) (actual time=1.023..1.024 rows=7 loops=1)
12          Buckets: 1024 Batches: 1 Memory Usage: 9kB
13          -> Seq Scan on publisher pu (cost=0.00..11.99 rows=6 width=17) (actual time=0.016..1.016 rows=7 loops=1)
14            Filter: ((publisher_name)::text ~~* ANY ('{%Nintendo%,%Sony%,%Microsoft%}'::text[]))
15            Rows Removed by Filter: 574
16          -> Index Scan using sales_game_id_key on sales s (cost=0.29..0.35 rows=1 width=9) (actual time=0.001..0.001 rows=1 loops=1)
17            Index Cond: (game_id = g.game_id)
18 Planning Time: 0.840 ms
19 Execution Time: 6.428 ms
```

Total rows: 19 | Query complete 00:00:00.044

QUERY 4: DECLINING PLATFORM POPULARITY OVER YEARS

```
59 /* QUERY 4: DECLINING PLATFORM POPULARITY OVER YEARS */
60
61 EXPLAIN ANALYZE
62 WITH yearly_counts AS (
63   SELECT
64     p.platform_name,
65     g.year_of_release,
66     COUNT(*) AS num_games
67   FROM game g
68   JOIN platform p ON g.platform_id = p.platform_id
69   WHERE g.year_of_release IS NOT NULL
70   GROUP BY p.platform_name, g.year_of_release
71 ), trend AS (
72   SELECT
73     *,
74     LAG(num_games) OVER (PARTITION BY platform_name ORDER BY year_of_release) AS prev_year
75   FROM yearly_counts
76 )
77   SELECT *
78   FROM trend
79   WHERE prev_year IS NOT NULL AND num_games < prev_year
80   ORDER BY platform_name, year_of_release;
```

Output:

The screenshot shows the PostgreSQL Explain Plan window. The top bar includes tabs for Data Output, Messages, and Notifications, along with various icons for file operations and SQL entry. The main area displays the query plan for a specific query, numbered from 1 to 18. The plan details the execution steps, including Subquery Scans, WindowAgg, GroupAggregate, Hash Joins, Seq Scans, and Sort operations. It also provides memory usage information like 'Memory: 1305kB' and 'Buckets: 4096 Batches: 1 Memory Usage: 34kB'. The bottom status bar indicates 'Total rows: 18' and 'Query complete 00:00:00.053'.

Step	Operation	Cost	Rows	Width	Actual Time
1	Subquery Scan on trend	(cost=1605.59..2438.34)	rows=5524	width=23	(actual time=11.612..13.854 loops=1)
2	Filter: ((trend.prev_year IS NOT NULL) AND (trend.num_games < trend.prev_year))				
3	Rows Removed by Filter: 142				
4	-> WindowAgg	(cost=1605.59..2230.16)	rows=16655	width=23	(actual time=11.602..13.835 loops=1)
5	-> GroupAggregate	(cost=1605.59..1938.69)	rows=16655	width=15	(actual time=11.589..13.736 loops=1)
6	Group Key: p.platform_name, g.year_of_release				
7	-> Sort	(cost=1605.59..1647.23)	rows=16655	width=7	(actual time=11.579..12.320 loops=1)
8	Sort Key: p.platform_name, g.year_of_release				
9	Sort Method: quicksort	Memory: 1305kB			
10	-> Hash Join	(cost=62.42..437.77)	rows=16655	width=7	(actual time=0.032..4.282 loops=1)
11	Hash Cond: (g.platform_id = p.platform_id)				
12	-> Seq Scan on game g	(cost=0.00..331.55)	rows=16655	width=8	(actual time=0.011..1.670 loops=16655)
13	Filter: (year_of_release IS NOT NULL)				
14	-> Hash	(cost=33.30..33.30)	rows=2330	width=7	(actual time=0.014..0.017 loops=31)
15	Buckets: 4096 Batches: 1	Memory Usage: 34kB			
16	-> Seq Scan on platform p	(cost=0.00..33.30)	rows=2330	width=7	(actual time=0.002..0.008 loops=31)
17	Planning Time: 0.267 ms				
18	Execution Time: 13.901 ms				

Total rows: 18 Query complete 00:00:00.053

QUERY 5: GAMES WHOSE SALES > AVG SALES OF THEIR GENRE

```
83  /* QUERY 5: GAMES WHOSE SALES > AVG SALES OF THEIR GENRE */
84
85  EXPLAIN ANALYZE
86  SELECT
87      g.name AS game_name,
88      ge.genre_name,
89      s.global_sales
90  FROM game g
91  JOIN genre ge ON g.genre_id = ge.genre_id
92  JOIN sales s ON g.game_id = s.game_id
93  WHERE s.global_sales > (
94      SELECT AVG(s2.global_sales)
95      FROM game g2
96      JOIN sales s2 ON g2.game_id = s2.game_id
97      WHERE g2.genre_id = g.genre_id
98  )
99  ORDER BY s.global_sales DESC;
```

Output:

Data Output		Messages	Notifications							
										SQL
	QUERY PLAN									
	text									
1	Sort (cost=1291.82..1305.70 rows=5552 width=37) (actual time=35178.035..35178.123 rows=3851 loops=1)									
2	Sort Key: s.global_sales DESC									
3	Sort Method: quicksort Memory: 339kB									
4	-> Hash Join (cost=556.64..946.51 rows=5552 width=37) (actual time=8.551..35174.366 rows=3851 loops=1)									
5	Hash Cond: (g.genre_id = ge.genre_id)									
6	-> Hash Join (cost=500.74..876.01 rows=5552 width=33) (actual time=8.519..35172.258 rows=3851 loops=1)									
7	Hash Cond: (g.game_id = s.game_id)									
8	Join Filter: (s.global_sales > (SubPlan 1))									
9	Rows Removed by Join Filter: 12804									
10	-> Seq Scan on game g (cost=0.00..331.55 rows=16655 width=32) (actual time=0.003..1.566 rows=16655 loops=1)									
11	-> Hash (cost=292.55..292.55 rows=16655 width=9) (actual time=4.394..4.394 rows=16655 loops=1)									
12	Buckets: 32768 Batches: 1 Memory Usage: 972kB									
13	-> Seq Scan on sales s (cost=0.00..292.55 rows=16655 width=9) (actual time=0.006..2.002 rows=16655 loops=1)									
14	SubPlan 1									
15	-> Aggregate (cost=728.69..728.70 rows=1 width=32) (actual time=2.110..2.110 rows=1 loops=16655)									
16	-> Hash Join (cost=389.20..725.48 rows=1281 width=5) (actual time=0.660..2.035 rows=1801 loops=16655)									
17	Hash Cond: (s2.game_id = g2.game_id)									
18	-> Seq Scan on sales s2 (cost=0.00..292.55 rows=16655 width=9) (actual time=0.001..0.593 rows=16655 loops=16...									
19	-> Hash (cost=373.19..373.19 rows=1281 width=4) (actual time=0.656..0.656 rows=1801 loops=16655)									
20	Buckets: 4096 (originally 2048) Batches: 1 (originally 1) Memory Usage: 150kB									
21	-> Seq Scan on game g2 (cost=0.00..373.19 rows=1281 width=4) (actual time=0.001..0.573 rows=1801 loops=1...									
22	Filter: (genre_id = g.genre_id)									
23	Rows Removed by Filter: 14854									
24	-> Hash (cost=30.40..30.40 rows=2040 width=12) (actual time=0.024..0.024 rows=13 loops=1)									
25	Buckets: 2048 Batches: 1 Memory Usage: 17kB									
26	-> Seq Scan on genre ge (cost=0.00..30.40 rows=2040 width=12) (actual time=0.015..0.016 rows=13 loops=1)									
27	Planning Time: 0.526 ms									
28	Execution Time: 35178.248 ms									
Total rows: 28		Query complete 00:00:35.385								

Conclusion

Query 5 was clearly the most expensive and met all characteristics of a high-complexity analytical query:

- Correlated subquery
- Multiple sequential scans
- Nested hash joins
- Sorting on large datasets
- $O(n^2)$ execution pattern across tables

Therefore, Query 5 was selected for performance tuning.

3. The Selected Query:

Query 5: GAMES WHOSE SALES > AVG SALES OF THEIR GENRE

Query Goal:

Return games whose global sales exceed the average global sales of their genre.

SQL:

```
83  /* QUERY 5: GAMES WHOSE SALES > AVG SALES OF THEIR GENRE */
84
85  EXPLAIN ANALYZE
86  SELECT
87      g.name AS game_name,
88      ge.genre_name,
89      s.global_sales
90  FROM game g
91  JOIN genre ge ON g.genre_id = ge.genre_id
92  JOIN sales s ON g.game_id = s.game_id
93  WHERE s.global_sales > (
94      SELECT AVG(s2.global_sales)
95      FROM game g2
96      JOIN sales s2 ON g2.game_id = s2.game_id
97      WHERE g2.genre_id = g.genre_id
98  )
99  ORDER BY s.global_sales DESC;
```

Why It Was Slow

- The correlated subquery runs once for every game row.
- Entire tables (game, sales) were scanned repeatedly.
- Lacked appropriate indexes on join/filter columns.
- Sorting by global_sales requires additional work after joins.

4. EXPLAIN ANALYZE BEFORE Indexing

Execution Time: 35,178.248 ms

Plan Characteristics:

- Multiple Seq Scans on:
 - game
 - sales
 - subquery tables
- Large hash joins repeatedly executed
- Sorting on 16,655+ rows

- Extensive row filtering during correlated comparison

This confirmed the query was performing large amounts of unnecessary work.

Data Output		Messages	Notifications					
QUERY PLAN								
text								
1	Sort (cost=1291.82..1305.70 rows=5552 width=37) (actual time=35178.035..35178.123 rows=3851 loops=1)							
2	Sort Key: s.global_sales DESC							
3	Sort Method: quicksort Memory: 339kB							
4	-> Hash Join (cost=556.64..946.51 rows=5552 width=37) (actual time=8.551..35174.366 rows=3851 loops=1)							
5	Hash Cond: (g.genre_id = ge.genre_id)							
6	-> Hash Join (cost=500.74..876.01 rows=5552 width=33) (actual time=8.519..35172.258 rows=3851 loops=1)							
7	Hash Cond: (g.game_id = s.game_id)							
8	Join Filter: (s.global_sales > (SubPlan 1))							
9	Rows Removed by Join Filter: 12804							
10	-> Seq Scan on game g (cost=0.00..331.55 rows=16655 width=32) (actual time=0.003..1.566 rows=16655 loops=1)							
11	-> Hash (cost=292.55..292.55 rows=16655 width=9) (actual time=4.394..4.394 rows=16655 loops=1)							
12	Buckets: 32768 Batches: 1 Memory Usage: 972kB							
13	-> Seq Scan on sales s (cost=0.00..292.55 rows=16655 width=9) (actual time=0.006..2.002 rows=16655 loops=1)							
14	SubPlan 1							
15	-> Aggregate (cost=728.69..728.70 rows=1 width=32) (actual time=2.110..2.110 rows=1 loops=16655)							
16	-> Hash Join (cost=389.20..725.48 rows=1281 width=5) (actual time=0.660..2.035 rows=1801 loops=16655)							
17	Hash Cond: (s2.game_id = g2.game_id)							
18	-> Seq Scan on sales s2 (cost=0.00..292.55 rows=16655 width=9) (actual time=0.001..0.593 rows=16655 loops=16...)							
19	-> Hash (cost=373.19..373.19 rows=1281 width=4) (actual time=0.656..0.656 rows=1801 loops=16655)							
20	Buckets: 4096 (originally 2048) Batches: 1 (originally 1) Memory Usage: 150kB							
21	-> Seq Scan on game g2 (cost=0.00..373.19 rows=1281 width=4) (actual time=0.001..0.573 rows=1801 loops=1...							
22	Filter: (genre_id = g.genre_id)							
23	Rows Removed by Filter: 14854							
24	-> Hash (cost=30.40..30.40 rows=2040 width=12) (actual time=0.024..0.024 rows=13 loops=1)							
25	Buckets: 2048 Batches: 1 Memory Usage: 17kB							
26	-> Seq Scan on genre ge (cost=0.00..30.40 rows=2040 width=12) (actual time=0.015..0.016 rows=13 loops=1)							
27	Planning Time: 0.526 ms							
28	Execution Time: 35178.248 ms							
Total rows: 28		Query complete 00:00:35.385						

5. Indexing Strategy

To reduce full table scans and improve join/filter performance, the following indexes were created:

```
/* Improve JOIN performance: sales - game */
CREATE INDEX idx_sales_game_id
ON sales(game_id);

/* Improve JOIN + GROUP BY performance on genre */
CREATE INDEX idx_game_genre_id
ON game(genre_id);

/* Improve sorting + filtering + join on global_sales */
CREATE INDEX idx_sales_global_sales_game_id
ON sales(global_sales, game_id);

/* Multi-column index to speed correlated subquery */
CREATE INDEX idx_game_genre_game_id
ON game(genre_id, game_id);
```

Rationale for Each Index

Index	Purpose
idx_sales_game_id	Accelerates joins between game-sales
idx_game_genre_id	Speeds up filtering/joins on genre_id
idx_sales_global_sales_game_id	Optimizes sorting/filtering on global_sales
idx_game_genre_game_id	Helps correlated subquery retrieve genre-matched rows faster

6. EXPLAIN ANALYZE AFTER Indexing

Execution Time: 27,158.193 ms

Performance Improvement:

- Improvement of ~8,020 ms
- ≈ 22.8% faster compared to original

Plan Improvements Noticed:

- Merge Join introduced (more efficient than Hash Join in this case)
- Index Scan replaced Seq Scan in several areas:
 - Index Scan using idx_game_genre_id on game
 - Index Only Scan using idx_game_genre_game_id on game g2
- Subquery now retrieves rows much more efficiently
- Overall reduced row processing during joins

Data Output		Messages	Notifications
		SQL	S
		QUERY PLAN	
text			
1		Sort (cost=1113.51..1127.39 rows=5552 width=37) (actual time=27157.943..27158.031 rows=3851 loops=1)	
2		Sort Key: s.global_sales DESC	
3		Sort Method: quicksort Memory: 339kB	
4		-> Hash Join (cost=556.54..768.21 rows=5552 width=37) (actual time=7.579..27156.523 rows=3851 loops=1)	
5		Hash Cond: (g.game_id = s.game_id)	
6		Join Filter: (s.global_sales > (SubPlan 1))	
7		Rows Removed by Join Filter: 12804	
8		-> Merge Join (cost=55.81..223.75 rows=16655 width=40) (actual time=0.045..12.400 rows=16655 loops=1)	
9		Merge Cond: (g.genre_id = ge.genre_id)	
10		-> Index Scan using idx_game_genre_id on game g (cost=0.29..948.62 rows=16655 width=32) (actual time=0.036..7.815 rows=16655 loops=1)	
11		-> Index Scan using genre_pkey on genre ge (cost=0.15..51.75 rows=2040 width=12) (actual time=0.002..0.010 rows=13 loops=1)	
12		-> Hash (cost=292.55..292.55 rows=16655 width=9) (actual time=4.180..4.181 rows=16655 loops=1)	
13		Buckets: 32768 Batches: 1 Memory Usage: 972kB	
14		-> Seq Scan on sales s (cost=0.00..292.55 rows=16655 width=9) (actual time=0.007..2.027 rows=16655 loops=1)	
15		SubPlan 1	
16		-> Aggregate (cost=394.20..394.21 rows=1 width=32) (actual time=1.629..1.629 rows=1 loops=16655)	
17		-> Hash Join (cost=54.72..391.00 rows=1281 width=5) (actual time=0.171..1.551 rows=1801 loops=16655)	
18		Hash Cond: (s2.game_id = g2.game_id)	
19		-> Seq Scan on sales s2 (cost=0.00..292.55 rows=16655 width=9) (actual time=0.000..0.598 rows=16655 loops=16655)	
20		-> Hash (cost=38.70..38.70 rows=1281 width=4) (actual time=0.168..0.168 rows=1801 loops=16655)	
21		Buckets: 4096 (originally 2048) Batches: 1 (originally 1) Memory Usage: 150kB	
22		-> Index Only Scan using idx_game_genre_game_id on game g2 (cost=0.29..38.70 rows=1281 width=4) (actual time=0.002..0.086 rows=1801 loops=16655)	
23		Index Cond: (genre_id = g.genre_id)	
24		Heap Fetches: 0	
25		Planning Time: 1.285 ms	
26		Execution Time: 27158.193 ms	

7. Why the Optimization Worked

Before:

- Database scanned entire game and sales tables multiple times.
- No efficient path to match rows by genre or game.
- Subquery forced repeated aggregations over large row sets.

After:

- Indexes allowed PostgreSQL to:
 - Perform index-only scans
 - Reduce join cost
 - Avoid full table scans
 - Reduce work inside subquery
- Query engine switched to Merge Join, which is more efficient for sorted/indexed data.

Even though correlated subqueries are inherently expensive, the applied indexes significantly reduced the computational load.

8. Conclusion

The optimization process successfully improved query performance by ~23%, demonstrating:

- Correct identification of the most expensive query
- Effective indexing strategy
- Clear, measurable performance gain
- Solid understanding of OLAP query optimization techniques

Github Repo Link : https://github.com/prajesh-1003/Video_Game_Sales_and_Ratings_Analytics